

Jacobian Calculations for nls()

Arkajyoti Bhattacharjee, Indian Institute of Technology, Kanpur
John C. Nash, University of Ottawa, Canada

26/05/2021

Contents

ISSUES	1
TODOS (mostly from nlsrc vignette nlsrc-devdoc.Rmd)	1
Jacobians in nls()	2
An example problem	2
Tools for Jacobians	3
numericDeriv() original version	3
numericDeriv() alternative pure-R version	8
Symbolic methods from nlsrc	11
numDeriv package	13
Comparisons	13
Performance results for different computing environments	15
Discussion of derivative computation for nonlinear least squares	15
Nomenclature	15
Numerical approximation near constraints	16
A case that blocks nls() from running	16
Appendix 1: Base R numericDeriv code	18
From nlsrc.R	18
From nlsrc.c	19
Appendix 2: numericDeriv() from nlsrc package (all in R)	21

ISSUES

- ExDerivs.R file causes a number of failures in the ORIGINAL numericDeriv.
- Need to verify nlsrc:: version of numericDeriv() matches all cases of nlspkg:: version
- Do we need to get a model frame? How? and How to use it?

TODOS (mostly from nlsrc vignette nlsrc-devdoc.Rmd)

- how to insert numerical derivatives when Deriv unable to get result

- approximations for `jacfn` beyond fwd approximation. How to specify??
- how to force numerical approximations in `nlfb()` in a manner consistent with that used in `optimx::optimr()`, that is, to surround the name of `jacfn` with quotes if it is a numerical approximation, or to provide a logical control to `nlxb()` for this purpose.

Jacobians in `nls()`

This document source is in file `DerivsNLS.Rmd`.

`nls()` needs Jacobians calculated at the current set of trial nonlinear model parameters to set up the Gauss-Newton equations. Unfortunately, `nls()` calls the Jacobian the “gradient,” and uses function `numericDerivs()` to compute them. This document is an attempt to describe different ways to compute the Jacobian for use in `nls()` and related software, and to evaluate the performance of these approaches.

In evaluating performance, we need to know the conditions under which the evaluation was conducted. Thus the computations included in this document, which is built using `Rmarkdown`, are specific to the computer in which the document is processed. We will add tables that give the results for different computing environments at the bottom.

An example problem

We will use the Hobbs weed infestation problem (Nash (1979), page 120).

```
# Data for Hobbs problem
ydat <- c(5.308, 7.24, 9.638, 12.866, 17.069, 23.192, 31.443,
          38.558, 50.156, 62.948, 75.995, 91.972) # for testing
tdat <- seq_along(ydat) # for testing

# A simple starting vector -- must have named parameters for nlxb, nls, wrapnlr.
start1 <- c(b1=1, b2=1, b3=1)
eunsc <- y ~ b1/(1+b2*exp(-b3*tt))
str(eunsc)

## Class 'formula' language y ~ b1/(1 + b2 * exp(-b3 * tt))
## .. attr(*, ".Environment")=<environment: R_GlobalEnv>

# Can we convert a string form of this "model" to a formula
ceunsc <- " y ~ b1/(1+b2*exp(-b3*tt))"
str(ceunsc)

## chr " y ~ b1/(1+b2*exp(-b3*tt))"

# Will be TRUE if we have made the conversion
print(as.formula(ceunsc)==eunsc)

## [1] TRUE

## LOCAL DATA IN DATA FRAMES
weeddata1 <- data.frame(y=ydat, tt=tdat)

## Put data in an Environment
weedenv <- list2env(weeddata1)
weedenv$b1 <- start1[[1]]
weedenv$b2 <- start1[[2]]
weedenv$b3 <- start1[[3]]
# Display content of the Environment
```

```

## Note that may need to do further commands to get everything
ls.str(weedenv)

## b1 :  num 1
## b2 :  num 1
## b3 :  num 1
## tt :  int [1:12] 1 2 3 4 5 6 7 8 9 10 ...
## y :  num [1:12] 5.31 7.24 9.64 12.87 17.07 ...

# Generate the residual "call"
rexpr<-call("-",eunsc[[3]], eunsc[[2]])
# Get the residuals
r0<-eval(rexpr, weedenv)
print(r0)

## [1] -4.576941 -6.359203 -8.685426 -11.883986 -16.075693 -22.194473
## [7] -30.443911 -37.558335 -49.156123 -61.948045 -74.995017 -90.972006
cat("Sumsquares at 1,1,1 is ",sum(r0^2),"\\n")

## Sumsquares at 1,1,1 is 23520.58

## Another way
ldata<-list2env(as.list(start1),envir=weedenv)
ldata

## <environment: 0x55ba9a62c658>
ls.str(ldata)

## b1 :  num 1
## b2 :  num 1
## b3 :  num 1
## tt :  int [1:12] 1 2 3 4 5 6 7 8 9 10 ...
## y :  num [1:12] 5.31 7.24 9.64 12.87 17.07 ...

eval(rexpr,envir=ldata)

## [1] -4.576941 -6.359203 -8.685426 -11.883986 -16.075693 -22.194473
## [7] -30.443911 -37.558335 -49.156123 -61.948045 -74.995017 -90.972006

## Do we need to get a model frame? How? and How to use it?

## Now ready to try things out.

```

Tools for Jacobians

numericDeriv() original version

numericDeriv is the R function used by nls() to evaluate Jacobians for its Gauss-Newton equations. The R source code is in the file nls.R. It calls a C function numeric_deriv in nls.c.

```

## Seems to work -- BUT note file ExDerivs.R has many "failures"??
theta <- c("b1", "b2", "b3")
ndeunsc<-nlspkg::numericDeriv(rexpr, theta, rho=weedenv)

## Registered S3 methods overwritten by 'nlspkg':
##   method      from
##   anova.nls    stats

```

```

## coef.nls      stats
## confint.nls   stats
## deviance.nls  stats
## df.residual.nls stats
## fitted.nls    stats
## formula.nls   stats
## logLik.nls    stats
## nobs.nls      stats
## plot.profile.nls stats
## predict.nls   stats
## print.nls     stats
## print.summary.nls stats
## profile.nls   stats
## residuals.nls stats
## summary.nls   stats
## vcov.nls      stats
## weights.nls   stats

print(ndeunsc)

## [1] -4.576941 -6.359203 -8.685426 -11.883986 -16.075693 -22.194473
## [7] -30.443911 -37.558335 -49.156123 -61.948045 -74.995017 -90.972006
## attr(,"gradient")
##      [,1]      [,2]      [,3]
## [1,] 0.7310585 -1.966119e-01 0.1966118813
## [2,] 0.8807971 -1.049936e-01 0.2099871635
## [3,] 0.9525741 -4.517674e-02 0.1355299950
## [4,] 0.9820137 -1.766276e-02 0.0706508160
## [5,] 0.9933071 -6.648064e-03 0.0332403183
## [6,] 0.9975274 -2.466440e-03 0.0147991180
## [7,] 0.9990890 -9.102821e-04 0.0063714981
## [8,] 0.9996643 -3.356934e-04 0.0026817322
## [9,] 0.9998765 -1.235008e-04 0.0011105537
## [10,] 0.9999547 -4.529953e-05 0.0004539490
## [11,] 0.9999828 -1.716614e-05 0.0001831055
## [12,] 0.9999943 -5.722046e-06 0.0000743866

print(sum(ndeunsc^2))

## [1] 23520.58

tndeunsc<-microbenchmark(ndeunsc<-nlspkg::numericDeriv(rexpr, theta, rho=weedenv))
print(tndeunsc)

## Unit: microseconds
##                                expr    min    lq
## ndeunsc <- nlspkg::numericDeriv(rexpr, theta, rho = weedenv) 9.936 10.3
##      mean median    uq    max neval
## 11.31204 10.506 10.73 38.828   100

## numericDeriv also has central difference option, as well as choice of eps parameter
## Central diff
ndeunsc2<-nlspkg::numericDeriv(rexpr, theta, rho=weedenv, central=TRUE)
print(ndeunsc2)

## [1] -4.576941 -6.359203 -8.685426 -11.883986 -16.075693 -22.194473
## [7] -30.443911 -37.558335 -49.156123 -61.948045 -74.995017 -90.972006

```

```

## attr("gradient")
##           [,1]           [,2]           [,3]
## [1,] 0.7310586 -1.966119e-01 1.966119e-01
## [2,] 0.8807971 -1.049936e-01 2.099872e-01
## [3,] 0.9525741 -4.517666e-02 1.355300e-01
## [4,] 0.9820138 -1.766271e-02 7.065082e-02
## [5,] 0.9933071 -6.648057e-03 3.324028e-02
## [6,] 0.9975274 -2.466509e-03 1.479906e-02
## [7,] 0.9990889 -9.102211e-04 6.371548e-03
## [8,] 0.9996647 -3.352378e-04 2.681902e-03
## [9,] 0.9998766 -1.233799e-04 1.110414e-03
## [10,] 0.9999546 -4.539623e-05 4.539581e-04
## [11,] 0.9999833 -1.670090e-05 1.837134e-04
## [12,] 0.9999939 -6.143885e-06 7.372897e-05

print(sum(ndeunsc2^2))

## [1] 23520.58

tndeunsc2<-microbenchmark(ndeunsc2<-nlspkg::numericDeriv(rexpr, theta, rho=weedenv, central=TRUE))
print(tndeunsc2)

## Unit: microseconds
##
##      ndeunsc2 <- nlspkg::numericDeriv(rexpr, theta, rho = weedenv,          expr
##      min      lq      mean median      uq      max neval      central = TRUE)
## 12.686 13.133 16.12278 13.473 13.7775 116.461   100

## Forward diff with smaller eps
ndeunscx<-nlspkg::numericDeriv(rexpr, theta, rho=weedenv, eps=1e-10)
print(ndeunscx)

## [1] -4.576941 -6.359203 -8.685426 -11.883986 -16.075693 -22.194473
## [7] -30.443911 -37.558335 -49.156123 -61.948045 -74.995017 -90.972006
## attr("gradient")
##           [,1]           [,2]           [,3]
## [1,] 0.7310597 -0.1966160568 0.1966071750
## [2,] 0.8807977 -0.1049915710 0.2099920238
## [3,] 0.9525714 -0.0451905180 0.1355182633
## [4,] 0.9820056 -0.0176747506 0.0706457115
## [5,] 0.9933032 -0.0066435746 0.0332534000
## [6,] 0.9975309 -0.0024513724 0.0148148160
## [7,] 0.9990941 -0.0009237056 0.0063593575
## [8,] 0.9996626 -0.0003552714 0.0026290081
## [9,] 0.9998757 -0.0001421085 0.0011368684
## [10,] 0.9999468 0.0000000000 0.0004973799
## [11,] 0.9998757 -0.0001421085 0.0001421085
## [12,] 1.0000178 0.0000000000 0.0001421085

print(sum(ndeunscx^2))

## [1] 23520.58

tndeunscx<-microbenchmark(ndeunscx<-nlspkg::numericDeriv(rexpr, theta, rho=weedenv, eps=1e-10))
print(tndeunscx)

## Unit: microseconds
##
##
## expr

```

```
## ndeunscx2 <- nls pkg::numericDeriv(rexpr, theta, rho = weedenv,      eps = 1e-10)
##      min      lq      mean median      uq      max neval
##  9.237  9.4625 10.18796  9.6255  9.921 40.619   100
```

```
## Central diff with smaller eps
```

```
ndeunscx2<-nls pkg::numericDeriv(rexpr, theta, rho=weedenv, central=TRUE, eps=1e-10)
print(ndeunscx2)
```

```
## [1] -4.576941 -6.359203 -8.685426 -11.883986 -16.075693 -22.194473
## [7] -30.443911 -37.558335 -49.156123 -61.948045 -74.995017 -90.972006
## attr("gradient")
##      [,1]      [,2]      [,3]
## [1,] 0.7310597 -1.966116e-01 1.966116e-01
## [2,] 0.8807977 -1.049916e-01 2.099876e-01
## [3,] 0.9525714 -4.518164e-02 1.355271e-01
## [4,] 0.9820145 -1.766587e-02 7.065459e-02
## [5,] 0.9933032 -6.643575e-03 3.325340e-02
## [6,] 0.9975309 -2.451372e-03 1.481482e-02
## [7,] 0.9990941 -9.059420e-04 6.359357e-03
## [8,] 0.9996981 -3.197442e-04 2.664535e-03
## [9,] 0.9998757 -1.421085e-04 1.136868e-03
## [10,] 0.9999468 -3.552714e-05 4.618528e-04
## [11,] 0.9999468 -7.105427e-05 2.131628e-04
## [12,] 1.0000178 0.000000e+00 7.105427e-05
```

```
print(sum(ndeunscx2^2))
```

```
## [1] 23520.58
```

```
tndeunscx2<-microbenchmark(ndeunscx2<-nls pkg::numericDeriv(rexpr, theta, rho=weedenv, central=TRUE, eps=1e-10),
print(tndeunscx2))
```

```
## Unit: microseconds
```

```
##
##      min      lq      mean median      uq      max neval      expr
##  11.985 12.359 13.27543 12.586 12.8045 54.864   100
##  11.985 12.359 13.27543 12.586 12.8045 54.864   100
```

```
## Add dir parameter
```

```
## to forward diff
```

```
ndeunsc<-nls pkg::numericDeriv(rexpr, theta, rho=weedenv)
print(ndeunsc)
```

```
## [1] -4.576941 -6.359203 -8.685426 -11.883986 -16.075693 -22.194473
## [7] -30.443911 -37.558335 -49.156123 -61.948045 -74.995017 -90.972006
## attr("gradient")
##      [,1]      [,2]      [,3]
## [1,] 0.7310585 -1.966119e-01 0.1966118813
## [2,] 0.8807971 -1.049936e-01 0.2099871635
## [3,] 0.9525741 -4.517674e-02 0.1355299950
## [4,] 0.9820137 -1.766276e-02 0.0706508160
## [5,] 0.9933071 -6.648064e-03 0.0332403183
## [6,] 0.9975274 -2.466440e-03 0.0147991180
## [7,] 0.9990890 -9.102821e-04 0.0063714981
## [8,] 0.9996643 -3.356934e-04 0.0026817322
## [9,] 0.9998765 -1.235008e-04 0.0011105537
## [10,] 0.9999547 -4.529953e-05 0.0004539490
```

```
## [11,] 0.9999828 -1.716614e-05 0.0001831055
## [12,] 0.9999943 -5.722046e-06 0.0000743866

ndeunscd<-nlspkg::numericDeriv(rexpr, theta, rho=weedenv, dir=-1)
print(ndeunscd)

## [1] -4.576941 -6.359203 -8.685426 -11.883986 -16.075693 -22.194473
## [7] -30.443911 -37.558335 -49.156123 -61.948045 -74.995017 -90.972006
## attr("gradient")
##      [,1]      [,2]      [,3]
## [1,] 0.7310586 -1.966119e-01 1.966119e-01
## [2,] 0.8807971 -1.049936e-01 2.099872e-01
## [3,] 0.9525741 -4.517663e-02 1.355300e-01
## [4,] 0.9820138 -1.766264e-02 7.065094e-02
## [5,] 0.9933071 -6.648064e-03 3.324032e-02
## [6,] 0.9975274 -2.466440e-03 1.479912e-02
## [7,] 0.9990890 -9.102821e-04 6.371498e-03
## [8,] 0.9996648 -3.352165e-04 2.682209e-03
## [9,] 0.9998765 -1.235008e-04 1.110554e-03
## [10,] 0.9999547 -4.529953e-05 4.539490e-04
## [11,] 0.9999838 -1.621246e-05 1.840591e-04
## [12,] 0.9999933 -6.675720e-06 7.343292e-05

ndeunscd<-nlspkg::numericDeriv(rexpr, theta, rho=weedenv, central=TRUE)
print(ndeunscd)

## [1] -4.576941 -6.359203 -8.685426 -11.883986 -16.075693 -22.194473
## [7] -30.443911 -37.558335 -49.156123 -61.948045 -74.995017 -90.972006
## attr("gradient")
##      [,1]      [,2]      [,3]
## [1,] 0.7310586 -1.966119e-01 1.966119e-01
## [2,] 0.8807971 -1.049936e-01 2.099872e-01
## [3,] 0.9525741 -4.517666e-02 1.355300e-01
## [4,] 0.9820138 -1.766271e-02 7.065082e-02
## [5,] 0.9933071 -6.648057e-03 3.324028e-02
## [6,] 0.9975274 -2.466509e-03 1.479906e-02
## [7,] 0.9990889 -9.102211e-04 6.371548e-03
## [8,] 0.9996647 -3.352378e-04 2.681902e-03
## [9,] 0.9998766 -1.233799e-04 1.110414e-03
## [10,] 0.9999546 -4.539623e-05 4.539581e-04
## [11,] 0.9999833 -1.670090e-05 1.837134e-04
## [12,] 0.9999939 -6.143885e-06 7.372897e-05

ndeunscd<-nlspkg::numericDeriv(rexpr, theta, rho=weedenv, central=TRUE, dir=-1)
print(ndeunscd)

## [1] -4.576941 -6.359203 -8.685426 -11.883986 -16.075693 -22.194473
## [7] -30.443911 -37.558335 -49.156123 -61.948045 -74.995017 -90.972006
## attr("gradient")
##      [,1]      [,2]      [,3]
## [1,] 0.7310586 -1.966119e-01 1.966119e-01
## [2,] 0.8807971 -1.049936e-01 2.099872e-01
## [3,] 0.9525741 -4.517666e-02 1.355300e-01
## [4,] 0.9820138 -1.766271e-02 7.065082e-02
## [5,] 0.9933071 -6.648057e-03 3.324028e-02
## [6,] 0.9975274 -2.466509e-03 1.479906e-02
```

```
## [7,] 0.9990889 -9.102211e-04 6.371548e-03
## [8,] 0.9996647 -3.352378e-04 2.681902e-03
## [9,] 0.9998766 -1.233799e-04 1.110414e-03
## [10,] 0.9999546 -4.539623e-05 4.539581e-04
## [11,] 0.9999833 -1.670090e-05 1.837134e-04
## [12,] 0.9999939 -6.143885e-06 7.372897e-05
```

numericDeriv() alternative pure-R version

This version (see Appendix 2) has C code replaced with R equivalents.

Try ExDerivs.R ??

```
andeunsc<-nlsalt::numericDeriv(rexpr, theta, rho=weedenv)
```

```
## Registered S3 methods overwritten by 'nlsalt':
```

```
## method      from
## anova.nls     nlspkg
## coef.nls     nlspkg
## confint.nls  MASS
## deviance.nls nlspkg
## df.residual.nls nlspkg
## fitted.nls   nlspkg
## formula.nls  nlspkg
## logLik.nls   nlspkg
## nobs.nls     nlspkg
## plot.profile.nls nlspkg
## predict.nls  nlspkg
## print.nls    nlspkg
## print.summary.nls nlspkg
## profile.nls  nlspkg
## residuals.nls nlspkg
## summary.nls  nlspkg
## vcov.nls     nlspkg
## weights.nls  nlspkg
```

```
print(andeunsc)
```

```
## [1] -4.576941 -6.359203 -8.685426 -11.883986 -16.075693 -22.194473
## [7] -30.443911 -37.558335 -49.156123 -61.948045 -74.995017 -90.972006
## attr(,"gradient")
##      [,1]      [,2]      [,3]
## [1,] 0.7310585 -1.966119e-01 0.1966118813
## [2,] 0.8807971 -1.049936e-01 0.2099871635
## [3,] 0.9525741 -4.517674e-02 0.1355299950
## [4,] 0.9820137 -1.766276e-02 0.0706508160
## [5,] 0.9933071 -6.648064e-03 0.0332403183
## [6,] 0.9975274 -2.466440e-03 0.0147991180
## [7,] 0.9990890 -9.102821e-04 0.0063714981
## [8,] 0.9996643 -3.356934e-04 0.0026817322
## [9,] 0.9998765 -1.235008e-04 0.0011105537
## [10,] 0.9999547 -4.529953e-05 0.0004539490
## [11,] 0.9999828 -1.716614e-05 0.0001831055
## [12,] 0.9999943 -5.722046e-06 0.0000743866
```

```
print(sum(andeunsc^2))
```



```
## [1] 23520.58
tandeunsc<-microbenchmark(andeunsc<-nlsalt::numericDeriv(rexpr, theta, rho=weedenv))
print(tandeunsc)

## Unit: microseconds
##                                expr      min      lq
## andeunsc <- nlsalt::numericDeriv(rexpr, theta, rho = weedenv) 31.191 31.932
##      mean median      uq      max neval
## 36.85921 32.3525 33.143 96.249   100

## numericDeriv also has central difference option, as well as choice of eps parameter
## Central diff
andeunsc2<-nlsalt::numericDeriv(rexpr, theta, rho=weedenv, central=TRUE)
print(andeunsc2)

## [1] -4.576941 -6.359203 -8.685426 -11.883986 -16.075693 -22.194473
## [7] -30.443911 -37.558335 -49.156123 -61.948045 -74.995017 -90.972006
## attr(,"gradient")
##           [,1]           [,2]           [,3]
## [1,] 0.7310586 -1.966119e-01 1.966119e-01
## [2,] 0.8807971 -1.049936e-01 2.099872e-01
## [3,] 0.9525741 -4.517666e-02 1.355300e-01
## [4,] 0.9820138 -1.766271e-02 7.065082e-02
## [5,] 0.9933071 -6.648057e-03 3.324028e-02
## [6,] 0.9975274 -2.466509e-03 1.479906e-02
## [7,] 0.9990889 -9.102211e-04 6.371548e-03
## [8,] 0.9996647 -3.352378e-04 2.681902e-03
## [9,] 0.9998766 -1.233799e-04 1.110414e-03
## [10,] 0.9999546 -4.539623e-05 4.539581e-04
## [11,] 0.9999833 -1.670090e-05 1.837134e-04
## [12,] 0.9999939 -6.143885e-06 7.372897e-05
print(sum(andeunsc2^2))

## [1] 23520.58
tandeunsc2<-microbenchmark(andeunsc2<-nlsalt::numericDeriv(rexpr, theta, rho=weedenv, central=TRUE))
print(tandeunsc2)

## Unit: microseconds
##                                expr
## andeunsc2 <- nlsalt::numericDeriv(rexpr, theta, rho = weedenv, central = TRUE)
##      min      lq      mean median      uq      max neval
## 40.734 41.7325 50.83806 42.679 52.775 107.232   100

## Forward diff with smaller eps
andeunscx<-nlsalt::numericDeriv(rexpr, theta, rho=weedenv, eps=1e-10)
print(andeunscx)

## [1] -4.576941 -6.359203 -8.685426 -11.883986 -16.075693 -22.194473
## [7] -30.443911 -37.558335 -49.156123 -61.948045 -74.995017 -90.972006
## attr(,"gradient")
##           [,1]           [,2]           [,3]
## [1,] 0.7310597 -0.1966160568 0.1966071750
## [2,] 0.8807977 -0.1049915710 0.2099920238
## [3,] 0.9525714 -0.0451905180 0.1355182633
## [4,] 0.9820056 -0.0176747506 0.0706457115
```

```
## [5,] 0.9933032 -0.0066435746 0.0332534000
## [6,] 0.9975309 -0.0024513724 0.0148148160
## [7,] 0.9990941 -0.0009237056 0.0063593575
## [8,] 0.9996626 -0.0003552714 0.0026290081
## [9,] 0.9998757 -0.0001421085 0.0011368684
## [10,] 0.9999468 0.0000000000 0.0004973799
## [11,] 0.9998757 -0.0001421085 0.0001421085
## [12,] 1.0000178 0.0000000000 0.0001421085

print(sum(andeunscx^2))

## [1] 23520.58

tandeunscx<-microbenchmark(andeunscx2<-nlsalt::numericDeriv(rexpr, theta, rho=weedenv, eps=1e-10))
print(tandeunscx)

## Unit: microseconds
##
## andeunscx2 <- nlsalt::numericDeriv(rexpr, theta, rho = weedenv,          expr
## min      lq      mean median      uq      max neval
## 30 30.6325 39.04353 31.245 53.5855 77.867 100

## Central diff with smaller eps
andeunscx2<-nlsalt::numericDeriv(rexpr, theta, rho=weedenv, central=TRUE, eps=1e-10)
print(andeunscx2)

## [1] -4.576941 -6.359203 -8.685426 -11.883986 -16.075693 -22.194473
## [7] -30.443911 -37.558335 -49.156123 -61.948045 -74.995017 -90.972006
## attr(,"gradient")
##          [,1]          [,2]          [,3]
## [1,] 0.7310597 -1.966116e-01 1.966116e-01
## [2,] 0.8807977 -1.049916e-01 2.099876e-01
## [3,] 0.9525714 -4.518164e-02 1.355271e-01
## [4,] 0.9820145 -1.766587e-02 7.065459e-02
## [5,] 0.9933032 -6.643575e-03 3.325340e-02
## [6,] 0.9975309 -2.451372e-03 1.481482e-02
## [7,] 0.9990941 -9.059420e-04 6.359357e-03
## [8,] 0.9996981 -3.197442e-04 2.664535e-03
## [9,] 0.9998757 -1.421085e-04 1.136868e-03
## [10,] 0.9999468 -3.552714e-05 4.618528e-04
## [11,] 0.9999468 -7.105427e-05 2.131628e-04
## [12,] 1.0000178 0.000000e+00 7.105427e-05

print(sum(andeunscx2^2))

## [1] 23520.58

tandeunscx2<-microbenchmark(andeunscx2<-nlsalt::numericDeriv(rexpr, theta, rho=weedenv, central=TRUE, eps=1e-10))
print(tandeunscx2)

## Unit: microseconds
##
## andeunscx2 <- nlsalt::numericDeriv(rexpr, theta, rho = weedenv,          central = TRUE, eps = 1e-10)
## min      lq      mean median      uq      max neval
## 39.091 40.172 43.81225 40.5885 41.647 116.849 100
```

The `dir` parameter allows us to use a backward difference for the derivative. This appears in `nlsModel()` for the case where a parameter is on an upper bound for the case `algorithm="port"`. It does not check for

nearness to the bound, and for the lower bound assumes that we are stepping AWAY from the bound in the default direction (`dir=+1`). None of the code addresses the issue where bounds are closer together than the step used for the finite difference, so there are situations where we could crash the code. Nor does the code check if the central difference is specified when near a bound.

- In the case of lower bounds, a central difference can overstep the bound when a parameter is “close” or on the bound.
- In the case of an upper bound, changing the `dir` will not change the derivative expression and steps in both forward and backward directions of the parameter are taken.

Symbolic methods from `nlsr`

The package `nlsr` has a function `model2rjfun()` that converts an expression describing how the residual functions are computed into an R function that computes the residuals at a particular set of parameters and sets the **attribute** “gradient” of the vector of residual values to the Jacobian at the particular set of parameters.

```
# nlsr has function model2rjfun. We can evaluate just the residuals
res0<-model2rjfun(eunsc, start1, data=weeddata1, jacobian=FALSE)
res0(start1)

## [1] -4.576941 -6.359203 -8.685426 -11.883986 -16.075693 -22.194473
## [7] -30.443911 -37.558335 -49.156123 -61.948045 -74.995017 -90.972006

# or the residuals and jacobian
## nlsr::model2rjfun forms a function with gradient (jacobian) attribute
funsc <- model2rjfun(eunsc, start1, data=weeddata1) # from nlsr: creates a function
tmodel2rjfun <- microbenchmark(model2rjfun(eunsc, start1, data=weeddata1))
print(tmodel2rjfun)

## Unit: microseconds
##                expr      min       lq      mean median
## model2rjfun(eunsc, start1, data = weeddata1) 83.915 85.399 100.8583 87.168
##                uq      max neval
## 100.5275 187.119   100

print(funsc)

## function(prm) {
##   if (is.null(names(prm)))
##     names(prm) <- names(pvec)
##   localdata <- list2env(as.list(prm), parent = data)
##   eval(residexpr, envir = localdata)
##   # Saves Jacobian matrix as "gradient" attribute (consistent with deriv())
## }
## <bytecode: 0x55ba9f4f5e80>
## <environment: 0x55ba9f438930>

print(funsc(start1))

## [1] -4.576941 -6.359203 -8.685426 -11.883986 -16.075693 -22.194473
## [7] -30.443911 -37.558335 -49.156123 -61.948045 -74.995017 -90.972006
## attr(,"gradient")
##          b1          b2          b3
## [1,] 0.7310586 -1.966119e-01 1.966119e-01
## [2,] 0.8807971 -1.049936e-01 2.099872e-01
## [3,] 0.9525741 -4.517666e-02 1.355300e-01
## [4,] 0.9820138 -1.766271e-02 7.065082e-02
```

```

## [5,] 0.9933071 -6.648057e-03 3.324028e-02
## [6,] 0.9975274 -2.466509e-03 1.479906e-02
## [7,] 0.9990889 -9.102212e-04 6.371548e-03
## [8,] 0.9996646 -3.352377e-04 2.681901e-03
## [9,] 0.9998766 -1.233793e-04 1.110414e-03
## [10,] 0.9999546 -4.539581e-05 4.539581e-04
## [11,] 0.9999833 -1.670114e-05 1.837126e-04
## [12,] 0.9999939 -6.144137e-06 7.372964e-05

print(environment(funsc))

## <environment: 0x55ba9f438930>

print(ls.str(environment(funsc)))

## data : <environment: 0x55ba9f431880>
## jacobian : logi TRUE
## modelformula : Class 'formula' language y ~ b1/(1 + b2 * exp(-b3 * tt))
## pvec : Named num [1:3] 1 1 1
## residexpr : expression({ .expr3 <- exp(-b3 * tt) .expr5 <- 1 + b2 * .expr3 .expr10 <- .expr5^2
## rjfun : function (prm)
## testresult : logi TRUE

print(ls(environment(funsc)$data))

## [1] "tt" "y"

eval(eunsc, environment(funsc))

## y ~ b1/(1 + b2 * exp(-b3 * tt))

vfunsc<-funsc(start1)
print(vfunsc)

## [1] -4.576941 -6.359203 -8.685426 -11.883986 -16.075693 -22.194473
## [7] -30.443911 -37.558335 -49.156123 -61.948045 -74.995017 -90.972006
## attr(,"gradient")
##          b1          b2          b3
## [1,] 0.7310586 -1.966119e-01 1.966119e-01
## [2,] 0.8807971 -1.049936e-01 2.099872e-01
## [3,] 0.9525741 -4.517666e-02 1.355300e-01
## [4,] 0.9820138 -1.766271e-02 7.065082e-02
## [5,] 0.9933071 -6.648057e-03 3.324028e-02
## [6,] 0.9975274 -2.466509e-03 1.479906e-02
## [7,] 0.9990889 -9.102212e-04 6.371548e-03
## [8,] 0.9996646 -3.352377e-04 2.681901e-03
## [9,] 0.9998766 -1.233793e-04 1.110414e-03
## [10,] 0.9999546 -4.539581e-05 4.539581e-04
## [11,] 0.9999833 -1.670114e-05 1.837126e-04
## [12,] 0.9999939 -6.144137e-06 7.372964e-05

tfunsc<-microbenchmark(funsc(start1))
print(tfunsc)

## Unit: microseconds
##          expr      min       lq      mean median        uq      max neval
## funsc(start1) 13.366 13.6985 14.78978 13.966 14.4215 54.57   100

```

numDeriv package

The package `numDeriv` includes a function `jacobian()` that acts on a user function `resid()` to produce the Jacobian at a set of parameters by several choices of approximation.

```
# We use the residual function (without gradient attribute) from nlsv
```

```
jeunsc<-jacobian(res0, start1)
jeunsc
```

```
##           [,1]           [,2]           [,3]
## [1,] 0.7310586 -1.966119e-01 1.966119e-01
## [2,] 0.8807971 -1.049936e-01 2.099872e-01
## [3,] 0.9525741 -4.517666e-02 1.355300e-01
## [4,] 0.9820138 -1.766271e-02 7.065082e-02
## [5,] 0.9933071 -6.648057e-03 3.324028e-02
## [6,] 0.9975274 -2.466509e-03 1.479906e-02
## [7,] 0.9990889 -9.102212e-04 6.371548e-03
## [8,] 0.9996647 -3.352378e-04 2.681902e-03
## [9,] 0.9998766 -1.233791e-04 1.110414e-03
## [10,] 0.9999546 -4.539572e-05 4.539580e-04
## [11,] 0.9999833 -1.670116e-05 1.837129e-04
## [12,] 0.9999939 -6.144205e-06 7.373002e-05
```

```
# Timings of the analytic jacobian calculations
```

```
tjeunsc<-microbenchmark(jeunsc<-jacobian(res0, start1))
print(tjeunsc)
```

```
## Unit: microseconds
```

```
##           expr      min       lq      mean median       uq
## jeunsc <- jacobian(res0, start1) 355.129 364.12 416.8117 374.04 419.2705
##           max neval
## 731.997    100
```

Note that the manual pages for `numDeriv` offer many options for the functions in the package. At 2021-5-27 we have yet to explore these.

Comparisons

In the following, we are comparing to `vfunsc`, which is the evaluated residual vector at `start1=c(1,1,1)` with “gradient” attribute (`jacobian`) included, as developed using package `nlsv`. This is taken as the “correct” result, even though it is possible that the generated order of calculations may introduce inaccuracies in the supposedly analytic derivatives.

`numericDeriv` computes a similar structure (residuals with “gradient” attribute): `ndeunsc`: the forward difference result with default `eps` ($1e-07$ according to manual) `ndeunsc2`: Central difference with default `eps` `ndeunscx`: Forward difference with smaller `eps=1e-10` `ndeunscx2`: Central difference with smaller `eps=1e-10`

`jeunsc`: `numDeriv::jacobian()` result with default settings.

```
## Matrix comparisons
```

```
attr(ndeunsc, "gradient")-attr(vfunsc, "gradient")
```

```
##           b1           b2           b3
## [1,] -4.066995e-08 -7.619266e-09 -5.198538e-08
## [2,] 1.016833e-08 3.631656e-09 -7.263312e-09
## [3,] 7.050552e-09 -8.473015e-08 1.577186e-08
## [4,] -8.764533e-08 -5.738229e-08 -8.889419e-09
```

```
## [5,] -3.542825e-08 -6.988878e-09 3.494439e-08
## [6,] -1.592723e-08 6.909055e-08 6.229383e-08
## [7,] 5.380365e-08 -6.095489e-08 -5.015294e-08
## [8,] -3.432289e-07 -4.556886e-07 -1.691883e-07
## [9,] -1.062480e-07 -1.214742e-07 1.395936e-07
## [10,] 9.833867e-08 9.627771e-08 -9.102750e-09
## [11,] -4.647158e-07 -4.649948e-07 -6.071033e-07
## [12,] 4.221287e-07 4.220910e-07 6.569545e-07
```

```
attr(ndeunsc2, "gradient")-attr(vfunsc,"gradient")
```

```
##           b1           b2           b3
## [1,] -5.513268e-11 1.371020e-11 5.962686e-11
## [2,] 6.850076e-13 -3.831403e-11 3.291006e-12
## [3,] -2.144829e-11 -1.208666e-10 6.925160e-11
## [4,] -2.665634e-11 4.123477e-11 -1.826495e-11
## [5,] 2.175706e-10 -7.825720e-11 9.793778e-11
## [6,] -2.416068e-10 -1.666460e-10 -1.735172e-10
## [7,] -8.350343e-11 5.415257e-11 -8.571976e-11
## [8,] 3.255913e-10 -9.864836e-11 2.024904e-10
## [9,] 1.379652e-11 -5.685668e-10 -1.631673e-10
## [10,] -9.296786e-11 -4.173983e-10 6.710748e-11
## [11,] -4.266865e-11 2.415372e-10 8.632699e-10
## [12,] -2.738492e-10 2.515432e-10 -6.717320e-10
```

```
attr(ndeunscx, "gradient")-attr(vfunsc,"gradient")
```

```
##           b1           b2           b3
## [1,] 1.078625e-06 -4.123528e-06 -4.758257e-06
## [2,] 5.790545e-07 2.014411e-06 4.852963e-06
## [3,] -2.771694e-06 -1.385826e-05 -1.171591e-05
## [4,] -8.202081e-06 -1.204434e-05 -5.113350e-06
## [5,] -3.931620e-06 4.482091e-06 1.311668e-05
## [6,] 3.569890e-06 1.513685e-05 1.576029e-05
## [7,] 5.191947e-06 -1.348438e-05 -1.219078e-05
## [8,] -2.074929e-06 -2.003370e-05 -5.289324e-05
## [9,] -8.676624e-07 -1.872920e-05 2.645423e-05
## [10,] -7.810096e-06 4.539581e-05 4.342184e-05
## [11,] -1.075608e-04 -1.254074e-04 -4.160402e-05
## [12,] 2.399048e-05 6.144137e-06 6.837890e-05
```

```
attr(ndeunscx2, "gradient")-attr(vfunsc,"gradient")
```

```
##           b1           b2           b3
## [1,] 1.078625e-06 3.173646e-07 -3.173646e-07
## [2,] 5.790545e-07 2.014411e-06 4.120706e-07
## [3,] -2.771694e-06 -4.976479e-06 -2.834131e-06
## [4,] 6.797035e-07 -3.162555e-06 3.768434e-06
## [5,] -3.931620e-06 4.482091e-06 1.311668e-05
## [6,] 3.569890e-06 1.513685e-05 1.576029e-05
## [7,] 5.191947e-06 4.279192e-06 -1.219078e-05
## [8,] 3.345221e-05 1.549344e-05 -1.736611e-05
## [9,] -8.676624e-07 -1.872920e-05 2.645423e-05
## [10,] -7.810096e-06 9.868671e-06 7.894701e-06
## [11,] -3.650654e-05 -5.435313e-05 2.945025e-05
## [12,] 2.399048e-05 6.144137e-06 -2.675369e-06
```

```
jeunsc-attr(vfunsc,"gradient")
```

```
##           b1           b2           b3
## [1,] -2.239464e-11  7.806283e-12 -1.156686e-12
## [2,] -2.267631e-12  2.974312e-11  2.957756e-11
## [3,] -8.948509e-12  3.630193e-11 -1.256267e-11
## [4,] -1.649125e-12  1.182179e-13  6.369841e-11
## [5,] -4.272493e-11 -1.109757e-11  3.501116e-11
## [6,]  1.867381e-10  1.793287e-11  3.319552e-11
## [7,]  1.090728e-11  1.840947e-11  9.946462e-12
## [8,]  2.035664e-10 -1.520996e-10  1.911593e-10
## [9,] -3.582228e-10  2.039028e-10 -1.905254e-10
## [10,] 3.202474e-10  8.291927e-11 -6.263366e-11
## [11,] 5.931922e-12 -1.779933e-11  3.682277e-10
## [12,] 4.132902e-10 -6.831839e-11  3.817154e-10
```

```
## Summary comparisons
```

```
max(abs(attr(ndeunsc, "gradient")-attr(vfunsc,"gradient")))
```

```
## [1] 6.569545e-07
```

```
max(abs(attr(ndeunsc2, "gradient")-attr(vfunsc,"gradient")))
```

```
## [1] 8.632699e-10
```

```
max(abs(attr(ndeunscx, "gradient")-attr(vfunsc,"gradient")))
```

```
## [1] 0.0001254074
```

```
max(abs(attr(ndeunscx2, "gradient")-attr(vfunsc,"gradient")))
```

```
## [1] 5.435313e-05
```

```
max(abs(jeunsc-attr(vfunsc,"gradient")))
```

```
## [1] 4.132902e-10
```

Performance results for different computing environments

Here we present tables of the results, preceded by identified descriptions of the machines we used. We use ideas and functions from the document `MachineSummary` to provide a characterization and identity for each machine used.

M21-LM20.1

?? still to be run

?? What machines provide a range of possibilities.

Discussion of derivative computation for nonlinear least squares

In no particular order, we comment on some issues relating to the Jacobian calculations in nonlinear least squares.

Nomenclature

R is not in step with many other areas of numerical computation when labelling different objects in the nonlinear least squares problem. In particular, R uses the term “gradient” when the object of interest is

the Jacobian matrix. In that it is useful in performing iterations of the Gauss-Newton or related equations to have the Jacobian associated with the residuals, and the rows of the Jacobian matrix are “gradients” of the respective residuals, we can accept the attribute name “gradient” to select the required information. Moreover, as in package `nlsr` it is very useful to have the Jacobian matrix as an attribute of the residual vector, since the main solver function, in this case `nlsr::nlfb()`, can be called with the same input for the arguments `res` and `jac`. These are the functions required to compute the residual and the Jacobian, and using the same function for both is very convenient, but needs some way to return both the residual vector and Jacobian matrix in a coherent fashion.

Numerical approximation near constraints

As far as we are aware there is no software that implements a fully safeguarded system to compute numerical approximation of the Jacobian (or gradients in general optimization) near constraints. The same statement applies even in the case of the much simpler bounds constraints. Users have a perverse tendency to devise ways to foil our best efforts. For example, they may decide that a good way to specify fixed (i.e., masked) parameters that they do not want to vary during a particular calculation is to specify the lower and upper bound of a parameter at the same value. Later runs may want the parameter constraints relaxed.

In `nlsr::nlxb()`, users may, in fact, specify masked parameters this way. This is a case of “if you can’t beat them, join them,” but it does provide an easily understood way for users to fix values.

More tricky is dealing with constraints that are close together. Note that these may arise from, for example, two linear (planar) constraints that approach at a narrow angle. In the apex where these constraints intersect, we will have tight bounds on parameters. If the constraint is not one that is imposed by the nature of the residual or objective function, for example, a `log()` or square root near zero, then we can generally proceed and allow the derivative approximator to evaluate outside the constraints. Things are decidedly nastier if we do have inadmissible values.

The issue of constraints and the need for a step in parameter values for derivative approximations was one of the motivations for trying to find analytic derivatives in package `nlsr` and the continuing effort to bring them into other R tools.

A case that blocks `nls()` from running

The following example shows that even supplying “good” starting points does not allow a usable Jacobian to be computed by `numericDeriv()`. This example is also in the document [RefactoringNLS](#).

```
# File: badJlogmod.R
# A problem illustrating poor numeric Jacobian
form<-y ~ 10*a*(8*b-log(0.075*c*x)) # the model formula
# This model uses log near a small argument, which skirts the dangerous
# value of 0. The parameters a, b, c could all be 1 "safely" as a start.
x<-1:20 # define x
a<-1.01
b<-0.9
c<-0.95
y <- 10*a*(8*b-log(0.075*c*x))+0.2*runif(20) # compute a y
df<-data.frame(x=x, y=y)
# plot(x,y) # for information
st<-c(a=1, b=1,c=1) # set the "default" starting vector
n0<-try(nls(form, start=st, data=df)) # and watch the fun as this fails.
```

```
## Error in nlsModel(formula, mf, start, wts, scaleOffset = scOff, nDcentral = nDcntr) :
##   singular gradient matrix at initial parameter estimates
```

```
library(nlsr) # but this will work
n1<-nlxb(form, start=st, data=df)
```



```
n1
```

```
## nlsr object: x
## residual sumsquares = 0.058082 on 20 observations
## after 5 Jacobian and 6 function evaluations
## name      coeff      SE      tstat      pval      gradient      JSingval
## a          1.01074      NA      NA      NA      -1.626e-11      501.9
## b          0.949944      NA      NA      NA      1.379e-11      25.63
## c          1.41028      NA      NA      NA      -1.222e-12      3.942e-15
```

```
jmod<-model2rjfun(form, pvec=st,data=data.frame(x=x, y=y)) # extract the model
Jatst<-jmod(st) # compute this at the start from package nlsr
Jatst<-attr(Jatst,"gradient") # and extract the Jacobian
#
# Now try to compute Jacobian produced by nls()
theta <- c("a","b","c") # the parameter names
env<-environment(form) # We need the environment of the formula
eform<-eval(form, envir=env) # and the evaluated expression
jnlsatst<-numericDeriv(form[[3L]], theta=theta, rho=env)
Jnls<-attr(jnlsatst,"gradient")
Jnls # from nls()
```

```
##      [,1] [,2] [,3]
## [1,] 98.416 80.8 -10.632
## [2,] 91.484 80.8 -10.632
## [3,] 87.429 80.8 -10.632
## [4,] 84.553 80.8 -10.632
## [5,] 82.321 80.8 -10.632
## [6,] 80.498 80.8 -10.632
## [7,] 78.957 80.8 -10.632
## [8,] 77.621 80.8 -10.632
## [9,] 76.443 80.8 -10.632
## [10,] 75.390 80.8 -10.632
## [11,] 74.437 80.8 -10.632
## [12,] 73.567 80.8 -10.632
## [13,] 72.766 80.8 -10.632
## [14,] 72.025 80.8 -10.632
## [15,] 71.335 80.8 -10.632
## [16,] 70.690 80.8 -10.632
## [17,] 70.083 80.8 -10.632
## [18,] 69.512 80.8 -10.632
## [19,] 68.971 80.8 -10.632
## [20,] 68.458 80.8 -10.632
```

```
Jatst # from nlsr -- analytic derivative
```

```
##      a b c
## [1,] 105.903 80 -10
## [2,] 98.971 80 -10
## [3,] 94.917 80 -10
## [4,] 92.040 80 -10
## [5,] 89.808 80 -10
## [6,] 87.985 80 -10
## [7,] 86.444 80 -10
## [8,] 85.108 80 -10
```

```
## [9,] 83.930 80 -10
## [10,] 82.877 80 -10
## [11,] 81.924 80 -10
## [12,] 81.054 80 -10
## [13,] 80.253 80 -10
## [14,] 79.512 80 -10
## [15,] 78.822 80 -10
## [16,] 78.177 80 -10
## [17,] 77.571 80 -10
## [18,] 76.999 80 -10
## [19,] 76.458 80 -10
## [20,] 75.945 80 -10

svd(Jnls)$d

## [1] 5.0277e+02 2.5680e+01 2.7199e-06

svd(Jatst)$d

## [1] 5.2370e+02 2.4390e+01 1.3047e-15

# what if we try to get closer to the parameters?
n0b<-try(nls(form, start=c(a=1, b=1,c=1.3), data=data.frame(x=x, y=y), trace=TRUE))

## Error in nlsModel(formula, mf, start, wts, scaleOffset = scOff, nDcentral = nDcntr) :
## singular gradient matrix at initial parameter estimates

# or even start at the solution?
n0c<-try(nls(form, start=coef(n1), data=data.frame(x=x, y=y), trace=TRUE))

## Error in nlsModel(formula, mf, start, wts, scaleOffset = scOff, nDcentral = nDcntr) :
## singular gradient matrix at initial parameter estimates
```

Appendix 1: Base R numericDeriv code

This code is in two files, nls.R and nls.c and is extracted here.

From nls.R

```
numericDeriv <- function(expr, theta, rho = parent.frame(), dir = 1,
  eps = .Machine$double.eps ^ (1/if(central) 3 else 2), central = FALSE)
## Note: this expr must be set up as a call to work properly according to JN??
## ?? we set eps conditional on central. But central set AFTER eps. Is this OK.
{
  cat("numericDeriv-Alt\n")
  dir <- rep_len(dir, length(theta))
  stopifnot(is.finite(eps), eps > 0)
  rho1 <- new.env(FALSE, rho, 0)
  if (!is.character(theta) ) {stop("'theta' should be of type character")}
  if (is.null(rho)) {
    stop("use of NULL environment is defunct")
    # rho <- R_BaseEnv;
  } else {
    if(! is.environment(rho)) {stop("'rho' should be an environment")}
    # int nprot = 3;
  }
  if( ! ((length(dir) == length(theta) ) & (is.numeric(dir) ) ) )
    {stop("'dir' is not a numeric vector of the correct length") }
```

```

if(is.na(central)) { stop("'central' is NA, but must be TRUE or FALSE") }
res0 <- eval(expr, rho) # the base residuals. ?? C has a check for REAL ANS=res0
if (any(is.infinite(res0)) ) {stop("residuals cannot be evaluated at base point")}
## CHECK_FN_VAL(res, ans); ?? how to do this. Is it necessary?
nt <- length(theta) # number of parameters
mr <- length(res0) # number of residuals
JJ <- matrix(NA, nrow=mr, ncol=nt) # Initialize the Jacobian
for (j in 1:nt){
  origPar<-get(theta[j],rho)
  xx <- abs(origPar)
  delta <- if (xx == 0.0) {eps} else { xx*eps }
  ## JN: I prefer eps*(xx + eps) which is simpler ?? Should we suggest / use a control switch
  prmx<-origPar+delta*dir[j]
  assign(theta[j],prmx,rho)
  res1 <- eval(expr, rho) # new residuals (forward step)
  if (central) { # compute backward step resids for central diff
    prmb <- origPar - dir[j]*delta
    assign(theta[j], prmb, envir=rho) # may be able to make more efficient later??
    resb <- eval(expr, rho)
    JJ[, j] <- dir[j]*(res1-resb)/(2*delta) # vectorized
  } else { ## forward diff
    JJ[,j] <- dir[j]*(res1-res0)/delta
  } # end forward diff
} # end loop over the parameters
attr(res0, "gradient") <- JJ
return(res0)
}

```

From nls.c

```

#include <stdlib.h>
#include <string.h>
#include <math.h>
#include <float.h>
#include <R.h>
#include <Rinternals.h>
#include "nls.h"
#include "internals.h"

#ifdef MIN
#define MIN(a,b) (((a)<(b))?(a):(b))
#endif

/*
 * call to numeric_deriv from R -
 * .Call("numeric_deriv", expr, theta, rho, dir = 1., eps = .Machine$double.eps, central=FALSE)
 * Returns: ans
 */
SEXP
numeric_deriv(SEXP expr, SEXP theta, SEXP rho, SEXP dir, SEXP eps_, SEXP centr,
              SEXP rho1)
{
  if(!isString(theta))
    error(_("'theta' should be of type character"));

```

```

    if (isNull(rho)) {
        error(_("use of NULL environment is defunct"));
        rho = R_BaseEnv;
    } else
        if(!isEnvironment(rho))
            error(_("'rho' should be an environment"));
    int nprot = 3;
    if(TYPEOF(dir) != REALSXP) {
        PROTECT(dir = coerceVector(dir, REALSXP)); nprot++;
    }
    if(LENGTH(dir) != LENGTH(theta))
        error(_("'dir' is not a numeric vector of the correct length"));
    Rboolean central = asLogical(centr);
    if(central == NA_LOGICAL)
        error(_("'central' is NA, but must be TRUE or FALSE"));
    //  SEXP rho1 = PROTECT(R_NewEnv(rho, FALSE, 0));
    //  nprot++;
    SEXP
    pars = PROTECT(allocVector(VECSXP, LENGTH(theta))),
    ans = PROTECT(duplicate(eval(expr, rho1)));
    double *rDir = REAL(dir), *res = NULL; // -Wall
#define CHECK_FN_VAL(_r_, _ANS_) do { \
    if(!isReal(_ANS_)) { \
        SEXP temp = coerceVector(_ANS_, REALSXP); \
        UNPROTECT(1); /*: _ANS_ *must* have been the last PROTECT() ! */ \
        PROTECT(_ANS_ = temp); \
    } \
    _r_ = REAL(_ANS_); \
    for(int i = 0; i < LENGTH(_ANS_); i++) { \
        if (!R_FINITE(_r_[i])) \
            error(_("Missing value or an infinity produced when evaluating the model")); \
    } \
} while(0)

    CHECK_FN_VAL(res, ans);

    const void *vmax = vmaxget();
    int lengthTheta = 0;
    for(int i = 0; i < LENGTH(theta); i++) {
        const char *name = translateChar(STRING_ELT(theta, i));
        SEXP s_name = install(name);
        SEXP temp = findVar(s_name, rho1);
        if(isInteger(temp))
            error(_("variable '%s' is integer, not numeric"), name);
        if(!isReal(temp))
            error(_("variable '%s' is not numeric"), name);
        // We'll be modifying the variable, so need to make a copy PR#15849
        defineVar(s_name, temp = duplicate(temp), rho1);
        MARK_NOT_MUTABLE(temp);
        SET_VECTOR_ELT(pars, i, temp);
        lengthTheta += LENGTH(VECTOR_ELT(pars, i));
    }
    vmaxset(vmax);
    SEXP gradient = PROTECT(allocMatrix(REALSXP, LENGTH(ans), lengthTheta));

```

```

double *grad = REAL(gradient);
double eps = asReal(eps_); // was hardcoded sqrt(DOUBLE_EPS) { ~= 1.49e-08, typically}
for(int start = 0, i = 0; i < LENGTH(theta); i++) {
double *pars_i = REAL(VECTOR_ELT(pars, i));
for(int j = 0; j < LENGTH(VECTOR_ELT(pars, i)); j++, start += LENGTH(ans)) {
    double
    origPar = pars_i[j],
    xx = fabs(origPar),
    delta = (xx == 0) ? eps : xx*eps;
    pars_i[j] += rDir[i] * delta;
    SEXP ans_del = PROTECT(eval(expr, rho1));
    double *rDel = NULL;
    CHECK_FN_VAL(rDel, ans_del);
    if(central) {
    pars_i[j] = origPar - rDir[i] * delta;
    SEXP ans_de2 = PROTECT(eval(expr, rho1));
    double *rD2 = NULL;
    CHECK_FN_VAL(rD2, ans_de2);
    for(int k = 0; k < LENGTH(ans); k++) {
        grad[start + k] = rDir[i] * (rDel[k] - rD2[k])/(2 * delta);
    }
    } else { // forward difference (previously hardwired):
    for(int k = 0; k < LENGTH(ans); k++) {
        grad[start + k] = rDir[i] * (rDel[k] - res[k])/delta;
    }
    }
    UNPROTECT(central ? 2 : 1); // ansDel & possibly ans
    pars_i[j] = origPar;
}
}
setAttrib(ans, install("gradient"), gradient);
UNPROTECT(nprot);
return ans;
}

```

Appendix 2: numericDeriv() from nlsalt package (all in R)

```

# File src/library/stats/R/nlsnd.R
# Part of the modified R package, https://www.R-project.org
#
# Copyright (C) 2000-2020 The R Core Team
# Copyright (C) 1999-1999 Saikat DebRoy, Douglas M. Bates, Jose C. Pinheiro
# J C Nash 2021
#
# This program is free software; you can redistribute it and/or modify
# it under the terms of the GNU General Public License as published by
# the Free Software Foundation; either version 2 of the License, or
# (at your option) any later version.
#
# This program is distributed in the hope that it will be useful,
# but WITHOUT ANY WARRANTY; without even the implied warranty of
# MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
# GNU General Public License for more details.

```

```

#
# A copy of the GNU General Public License is available at
# https://www.R-project.org/Licenses/

####
####          numeric Jacobian for Nonlinear least squares for R
####

numericDeriv <- function(expr, theta, rho = parent.frame(), dir = 1,
                          eps = .Machine$double.eps ^ (1/if(central) 3 else 2), central = FALSE)
## Note: this expr must be set up as a call to work properly according to JN??
## ?? we set eps conditional on central. But central set AFTER eps. Is this OK.
{
  ndtrace<-FALSE
  #   if(ndtrace) cat("numericDeriv-Alt\n")
  dir <- rep_len(dir, length(theta))
  stopifnot(is.finite(eps), eps > 0)
  rho1 <- new.env(FALSE, rho, 0)
  if (!is.character(theta) ) {stop("'theta' should be of type character")}
  if (is.null(rho)) {
    stop("use of NULL environment is defunct")
    #       rho <- R_BaseEnv;
  } else {
    if(! is.environment(rho)) {stop("'rho' should be an environment")}
    #   int nprot = 3;
  }
  if( ! ((length(dir) == length(theta) ) & (is.numeric(dir) ) ) )
    {stop("'dir' is not a numeric vector of the correct length") }
  if(is.na(central)) { stop("'central' is NA, but must be TRUE or FALSE") }
  res0 <- eval(expr, rho) # the base residuals. ?? C has a check for REAL ANS=res0
  if (any(is.infinite(res0)) ) {stop("residuals cannot be evaluated at base point")}
  ## CHECK_FN_VAL(res, ans); ?? how to do this. Is it necessary?
  nt <- length(theta) # number of parameters
  mr <- length(res0) # number of residuals
  JJ <- matrix(NA, nrow=mr, ncol=nt) # Initialize the Jacobian
  for (j in 1:nt){
    origPar<-get(theta[j],rho)
    xx <- abs(origPar)
    delta <- if (xx == 0.0) {eps} else { xx*eps }
    ## JN: I prefer eps*(xx + eps)  which is simpler ?? Should we suggest / use a control switch
    prmx<-origPar+delta*dir[j]
    assign(theta[j],prmx,rho)
    res1 <- eval(expr, rho) # new residuals (forward step)
    if (central) { # compute backward step resids for central diff
      prmb <- origPar - dir[j]*delta
      assign(theta[j], prmb, envir=rho) # may be able to make more efficient later??
      resb <- eval(expr, rho)
      JJ[, j] <- dir[j]*(res1-resb)/(2*delta) # vectorized
    } else { ## forward diff
      JJ[,j] <- dir[j]*(res1-res0)/delta
    } # end forward diff
    assign(theta[j],origPar,rho) # restore the parameter value !! IMPORTANT
  } # end loop over the parameters
  attr(res0, "gradient") <- JJ
  if (ndtrace){

```

```

        cat("par:")
        for (j in 1:nt){ cat(get(theta[j],rho)," ") }
        cat("\n")
        print(res0)
    }
    return(res0)
}

```

Nash, John C. 1979. *Compact Numerical Methods for Computers : Linear Algebra and Function Minimisation*. Book. Hilger: Bristol.